

Structure Learning of Generative Models for Agency Phenotyping with Tensor Networks

Philip Wilson¹, Mahault Albarracin^{2,3}, Nicolás Hinrichs⁴, Jasmine Moore¹,
Daniel Polani⁵, Karl J. Friston⁶, and Axel Constant⁷

¹ Independent Researcher

² LANCI, Université du Québec à Montréal, Canada

³ Centre of Excellence for AI and Robotics, Sheffield Hallam University, UK

⁴ Max Planck Institute for Human Cognitive and Brain Sciences, Leipzig, Germany

⁵ University of Hertfordshire, Hatfield, UK

⁶ Wellcome Centre for Human Neuroimaging, UCL, London, UK

⁷ University of Sussex, Brighton, UK

Abstract. An agentic system, we argued in earlier work [1], should be intentional, rational, and explainable; active inference secures the three in a partially observable Markov decision process whose tensors are open to inspection, and empowerment then reads off how much control the agent has over what it will observe. That case was built on models we wrote by hand. Its open question was whether the model, and the agency-relevant structure inside it, could instead be recovered from behaviour. This paper answers it, and the answer turns out to be an entire interpretable model, not a single number. A matrix product state (an interpretable tensor network) is trained unsupervised on T-maze rollouts, and from the trained network we recover, each validated against the true environment: normalised hidden states, by a gauge-invariant criterion of predictive equivalence that earlier tensor-network models could not supply; the labelled likelihood $A = p(o \mid s)$ and transition $B = p(s' \mid s, a)$, matched to total-variation 0.000; the conditional-independence graph among observables, by exact mutual information; the factorisation of the observation, in which context is recovered as a near-independent modality; and the number of hidden states, fixed by model selection read as active-inference model expansion and Bayesian model reduction. A gauge transformation then makes the bond index itself a labelled state. Because the recovered model is fully interpretable, the controllability phenotype now reads off a learned model rather than a stipulated one. The construction assumes no active-inference agent, and so transfers to any system whose sensorimotor stream can be modelled.

Keywords: Active inference · Agency · Computational phenotyping · Empowerment · Tensor networks · Structure learning · World models

1 Introduction

Agency in artificial systems is asserted far more often than it is measured. We argued in earlier work [1] that agency worth the name is intentional, rational,

and explainable, that active inference [2] secures the three at once in a POMDP whose likelihood and transition tensors are open to inspection, and that empowerment [3,4], the channel capacity from an action to its next observation, measures the control an agent holds over what it will see. That argument ran on generative models we specified by hand. What it left undone was to obtain such a model, and the agency-relevant structure inside it, from behaviour.

That is the problem this paper solves, and it is a problem in *structure learning* of a generative model, not in learning “agency” directly. Agency sets the target: the structure worth recovering is exactly what the three criteria and the empowerment probe read off a world model, so that once the model is learned the phenotype runs on it. Phenotyping, taken from computational psychiatry [5,6], has two steps: fit a generative model to behaviour, then read behaviour off the fitted parameters. Earlier we did only the second, on models specified by hand [1]; here we do the first, and find that the fitted model can be made fully interpretable rather than left as a black-box density.

Concretely, we train a *matrix product state* (MPS), a one-dimensional tensor network, as an unsupervised generative model of T-maze rollouts, following Wauthier et al. [7]. Tensor networks are attractive here because they admit exact marginalisation: the conditionals the empowerment probe needs, and the mutual informations a dependency analysis needs, are available in closed form from reduced density matrices. But a trained density model is not yet a *structured* one. Wauthier et al. observed that although the bonds between tensors “can be regarded as internal states, they are not normalised and therefore not usable,” and consequently bypassed hidden-state recovery entirely. Our central technical contribution is to supply exactly that missing piece and to build a full structural readout on top of it.

Contributions. From MPS models trained on the minimal four-observation T-maze and the full `pymdp` T-maze (24 composite observations), we recover, and validate against analytic ground truth:

1. **Hidden states** by gauge-invariant fidelity clustering, corrected to a *predictive-equivalence* criterion that is robust to gauge coherences the model writes into the bond (Sec. 3);
2. **A labelled generative model**: emission $A = p(o \mid s)$ and transition $B = p(s' \mid s, a)$, matched to the environment up to a state permutation (Sec. 4);
3. the **conditional-independence graph** over observables, by exact mutual information (Sec. 5);
4. the **observation factorisation**, from subsystem mutual information of the single-site reduced density matrix (Sec. 6);
5. the **number of hidden states** by model selection, framed as active inference model expansion and Bayesian model reduction (Sec. 7); and
6. **symbolic states**: a gauge that makes the bond index itself a labelled hidden state (Sec. 8).

On the learned model we then phenotype control capacity (Sec. 9). All code accompanies the paper.




Fig. 1. The T-maze, reproduced from [1]. The arrow marks the observation revealing where the cheese is; entering an arm is an absorbing trap, so the epistemically optimal first move is the cue.

2 Background

We recap only what the pipeline needs; the three agency criteria, the empowerment probe, and the design of the T-maze are developed in full in our earlier paper [1], and the tensor-network machinery is new here.

2.1 The task and the analytic model

The minimal T-maze is a two-step POMDP with three actions (Left, Right, Cue) and four observations (Cheese, Shock, Left, Right); cheese and shock are placed left/right with probability $1/2$ and the cue reveals their location, while entering an arm is an absorbing trap (Fig. 1) [1]. The full `pymdp` [8] T-maze factors the observation into Position (4), Reward (3), and Context (2), yielding 24 composite observations over a three-step episode beginning with a null action. Both admit an exact generative model $q(o, a)$, which we use only as ground truth for validation—the learner never sees it.

2.2 What is learned, and from what

The training set is the space of action–observation trajectories the environment can produce, not the log of any one agent. We enumerate every action sequence and pair it with the observations the environment returns, weighting each trajectory by its true probability, so the action distribution is uniform and carries no policy. What the MPS fits is therefore the environment’s own input–output law, and the structure we later recover is the environment’s generative model (its hidden states, likelihood, and transition), which is exactly why it can be checked against the analytic $q(o, a)$. Learning instead from a *particular* agent’s rollouts would recover the model that agent’s policy induces, a different and agent-specific object; we return to this in Sec. 10. Empowerment is indifferent to the choice, as it already maximises over the action distribution.

2.3 MPS as a generative model

An MPS represents the amplitude of a length- L sequence as a chain of contracted tensors, one per time step, each carrying a physical index (here the pair

[action, observation]) and two bond indices. Trained by the Born rule $p(x) = |\psi(x)|^2/Z$ [9], it is a *Born machine*, and admits exact contraction of any marginal or conditional. Non-negative MPS are exactly hidden Markov models, with the bond index a classical hidden state [10]; a generic complex Born machine is strictly more expressive and has no such exact form, so hidden states must be recovered rather than read off (Sec. 3). We train each model with a DMRG-style cumulant gradient sweep and a learning-rate schedule that shrinks the step when the loss stalls [9]; the minimal model reaches $\|p - q\|_1 = 0.13$ and the full model $\|p - q\|_1 = 0.005$.

2.4 Empowerment

We borrow one quantity from that work [1], $\max_{p(a)} I(A; O')$, the largest mutual information an agent can drive between its next action and its next observation, evaluated by Blahut–Arimoto [4]. It distinguishes the value computed under the true environment (*objective*) from the value under the agent’s beliefs (*subjective*), and both from the *actual* value once preferences narrow the action distribution. Only one variant is empirical here: computed under the learned model, the subjective value approaches the objective one as the fit tightens. We report it over the full observation and over the reward modality alone.

3 Recovering hidden states

The bond of a trained MPS is defined only up to a gauge transformation $A^{(n)} \rightarrow A^{(n)}G$, $A^{(n+1)} \rightarrow G^{-1}A^{(n+1)}$, so the bare bond index has no intrinsic labelling, so any state we read off it must be invariant to that freedom. Each history h (a prefix of actions and observations) drives the network into a bond vector $|b_h\rangle$; the quantity $|\langle b_i | b_j \rangle|^2$ is gauge-invariant, and grouping histories by it gives a first partition into states (Fig. 2, left/centre).

On the minimal maze this yields exactly four states (cheese-trap, shock-trap, and one per cue context), merging the two arms of entry because position is unobservable there. On the full maze, however, raw fidelity clustering *over-separates*: the model writes the uninformative context bit of the mid-step observation into the bond as a coherence that does not affect the future, splitting otherwise-identical states. The robust criterion is *predictive equivalence*—the predictive-state / observable-operator definition of state [11,12]: two histories are the same state iff they induce the same future conditional $p(o' | a', h)$. Agglomerating the action-conditioned predictive signatures and merging any two whose futures agree to within an L1 tolerance (Fig. 2, right) collapses the gauge duplicates: the four context-bit pairs merge at ≈ 0 distance, the two start histories at 1.2, and the seven genuine states stay ≥ 2.0 apart, so any threshold in that gap returns them (unresolved start, four arm–outcome traps, two cue contexts). The count does not hang on the threshold: it is corroborated by the bond-dimension model selection of Sec. 7 and by the fact that all seven recovered emissions match




Fig. 2. Recovering states beyond bond-ray fidelity. Left: on the full maze, the fidelity clustering used in our earlier pipeline [1] over-separates into nine apparent blocks (twelve histories), because the model writes the uninformative context bit into the bond as a coherence that does not affect the future. Right: predictive equivalence merges histories whose predicted futures agree; the gauge duplicates collapse at ≈ 0 distance while the seven genuine states stay $\gtrsim 2$ apart, so a threshold in the gap recovers the correct count.

ground truth (Sec. 4). What the learned latent encodes is therefore distinguishability rather than raw history: whether the agent entered from the left or the right arm collapses into one state when position cannot be seen, and splits into two once it can.

4 A labelled generative model: emission and transition

Given the states, we reconstruct the factorised model. Because probabilities are Born-squared, we push amplitudes to probabilities *before* forming any stochastic object: attempting to read a stochastic matrix off the complex tensor, or to gauge-transform the model into a non-negative form, is either ill-defined or NP-hard for a genuine Born machine [10]. The **emission** $A = p(o | s)$ is estimated by conditioning the exact joint on each state’s histories and marginalising the action; the **transition** $B = p(s' | s, a)$ by re-contracting the bond after appending (a, o) , classifying the resulting state by predictive equivalence, and accumulating Born weight. States are matched to the environment by the Hungarian algorithm, the only residual gauge freedom being a discrete permutation.

The reconstruction matches ground truth to measurement precision (Fig. 3, Table 1). On the minimal maze the emission matches ground truth to total-variation 0.000 for all four states. On the full maze, all seven emissions match to mean total-variation 0.000: the four traps are deterministic, and the two




Fig. 3. Recovered generative model (full maze). Left: emission $A = p(r_3 | a_3, s)$ over seven states and four actions $\times$ three rewards; each state's pattern matches analytic ground truth (mean TV 0.000). Right: transition $B = p(s_2 | s_1 = \text{start}, a_2)$, recovering the exact 0.50/0.50 arm and cue branching.

cue states carry the correct *opposite* action-conditioned reward patterns (under Right, context 0 yields cheese and context 1 yields shock). The recovered transition reproduces the branching structure of the environment exactly: from the start, Right and Left each split 0.50/0.50 over their two arm–outcome states, Cue splits 0.50/0.50 over the two context states, and staying at the centre is deterministic. This is the object Wauthier et al. could not extract: a normalised, labelled, ground-truth-comparable A and B read from a network trained on behaviour alone.

5 The dependency graph

Because the learned joint is available in closed form, its dependency structure follows from exact mutual information. We form the pairwise MI matrix over the observed variables and take its maximum-weight spanning tree (Chow–Liu [13]), the tensor-network analogue of reading correlation structure off a quantum state [14] (Fig. 4). The recovered tree is the correct temporal chain: the action at each step is strongly coupled to its observation and the observation to the next. Conditional mutual information supplies exact independence certificates: $I(a_1; o_3 | o_2) = 0$ (the first action affects the last observation only through the intervening one) and, in the minimal maze, $I(a_1; o_3 | o_2) = 0$ with $I(o_2; o_3 | a_2) = 1.31$ bits (dependence where the model should be dependent). We validate against the *observable-projected* ground-truth graph, since marginalising the latent state induces real dependencies among observables that a naive comparison to the full dynamic Bayes net would misread as error.




Fig. 4. Dependency and factorisation (full maze). Left: pairwise mutual-information matrix over the six observed variables. Centre: the Chow-Liu tree recovers the temporal action-observation chain. Right: subsystem mutual information of the single-site reduced density matrix—the model discovers context as a near-independent factor ($I(\text{pos:ctx}), I(\text{rew:ctx}) \approx 0.05$ bit) while position and reward are coupled (1.0 bit), the causally correct structure.

6 Discovering the observation factorisation

The full-maze model is fed a *flat* 24-dimensional observation; the Position $\times$ Reward $\times$ Context factorisation is assumed only in analysis. We test whether the learned tensors respect it, from the single-site reduced density matrix $\rho = \text{Tr}_{\text{rest}} |\psi\rangle\langle\psi|$. Subsystem quantum mutual information $I(A:B) = S(\rho_A) + S(\rho_B) - S(\rho_{AB})$ vanishes iff the factors are independent [15]. The model discovers context as a *near-independent* modality ($I(\text{pos:ctx}), I(\text{rew:ctx}) \approx 0.05$ bits) while position and reward remain coupled (1.0 bit)—which is causally correct, since position gates whether a reward is possible (Fig. 4, right). A search over groupings of the 24 indices confirms it: the trace-distance residual $\frac{1}{2}\|\rho - \bigotimes_k \rho_k\|_1$ is far lower for (position, reward) $\times$ context (0.21) than for full separation (0.54). The model organises the observation along the environment’s causal axes without being told they exist.

7 Selecting the number of states

That the recovered dimensionality is a *model-selection* result, not an artefact of an SVD threshold, is shown by sweeping the capped bond dimension χ and training each model to convergence (Fig. 5). On the minimal maze the fit improves sharply until $\chi = 4$ (fit L1 $0.17 \rightarrow 0.03$) and then goes flat: extra bond dimension is unused “spare capacity.” We read this through active inference structure learning [16,17]: a large- χ model has spare capacity; training commits capacity to states that raise model evidence (accuracy gain exceeding complexity cost) and leaves the rest unused, so the selected count is the fixed point of a free-energy trade-off over model structure—model expansion followed by Bayesian model reduction. A description-length penalty is one state more conservative




Fig. 5. State-count selection (minimal maze). Fit L_1 (blue) drops $\sim 6\times$ from $\chi = 3$ to $\chi = 4$ and then plateaus; bond dimension beyond four is spare capacity. The description-length score (grey) is one state more conservative.

here, because the fourth state adds interpretable low-probability structure that the fit needs but the likelihood barely rewards; the L_1 knee is the decisive signal.

8 Symbolic states by gauge fixing

Finally we remove the post-hoc character of state extraction by choosing the residual gauge that rotates the bond into the cluster basis. From the k cluster representatives we build an orthonormal frame (the eigenbasis of the history-aggregated bond density is more robust than a raw QR) and complete it to a unitary U ; applying U as the gauge preserves the represented distribution exactly (here to 4×10^{-14}) while the bond index becomes a labelled state that the model *emits*, $p_i(h) = |\langle q_i | U^\dagger | b_h \rangle|^2$ (Fig. 6, right) [18,19]. On the minimal maze the emitted distributions agree with the fidelity clustering (adjusted Rand index 1.0), near one-hot for the trap states and softer for the cue states—the expected signature of a genuinely complex Born machine, for which an exact classical-state gauge does not exist [10], quantified rather than hidden.

9 Phenotyping control capacity on the learned model

With the model recovered, we phenotype it. Subjective potential empowerment computed from the converged full-maze network reproduces the analytic regimes to three decimals (Fig. 6, left; Table 1): 2.00 bits at the start, 0.00 in the trap, and 2.00 bits (full) / 1.585 bits (reward) post-cue, against analytic 2.00/1.585. Restricting the probe to the reward modality returns the three control regimes of the minimal maze (1.00/0.00/1.585 bits): reading empowerment only over outcomes the agent cares about isolates the control that matters for behaviour, which the full observation space dilutes. The subjective–objective gap is now




Fig. 6. Left: empowerment phenotyping on the learned full-maze model matches analytic across the zero-, intermediate-, and maximal-control regimes, full-observation and reward-only. Right: the gauge-fixed bond emits labelled states (minimal maze); trap histories are near one-hot, cue histories softer.

Table 1. Structure recovery against analytic ground truth. Emission/transition are total-variation distances; empowerment is bits (full / reward). “Learned” is the converged full-maze MPS ($\|p - q\|_1 = 0.005$).

Quantity	Analytic	Learned	Error
Emission $A = p(o s)$ (mean TV over 7 states)	—	—	0.000
Transition B branching ($\text{Right} \rightarrow \{\text{cheese}, \text{shock}\}$)	0.50/0.50	0.51/0.49	≤ 0.01
Empowerment, start	2.00/1.00	2.00/1.00	0.00
Empowerment, trap	0.00/0.00	0.00/0.00	0.00
Empowerment, post-cue	2.00/1.585	2.00/1.585	< 0.001
Hidden-state count (minimal / full)	4 / 7	4 / 7	exact

empirical and small: a bond-truncated checkpoint attenuates post-cue empowerment (0.65 bits reward-only), and it approaches the objective value as the fit improves. Acting to resolve uncertainty increases empowerment: visiting the cue lifts it from the intermediate to the maximal regime, though empowerment is no part of the training objective.

10 Discussion

What is new. Earlier tensor-network work learned $p(o, a)$ but could not expose normalised hidden states, and the phenotyping programme it was meant to serve was demonstrated only on hand-specified models. We supply the missing state construction, gauge-invariant and based on predictive equivalence, and build on it a complete structural account: a labelled likelihood and transition validated to $\sim 10^{-3}$, the conditional-independence graph, the observation factorisation, a principled state count, and a symbolic-state gauge. The controllability probe then runs on the learned model and matches the analytic regimes. The pipeline

is architecture-agnostic: the tensor network is not an active-inference agent, so recovering active-inference-style structure from it addresses the circularity worry that active inference might satisfy the agency criteria only by construction.

Senses of structure learning. We are precise about the phrase. The model performs unsupervised discovery of latent dimensionality (Sec. 7), of the task-relevant state abstraction (Sec. 3), of the dependency graph (Sec. 5), and of the observation factorisation (Sec. 6), and it yields a labelled factored model (Sec. 4). It does not search over network topology; the tensor-network structure-search literature [20,21] would be the route to that, and is a natural next step on richer environments.

What the agency claim is, and is not. The contribution to agency is the recovered model, not the empowerment number. The states, the likelihood, and the transition are the belief-like and world-model structures that the intentionality, rationality, and explainability criteria are read off; recovering them from behaviour is what lets those criteria reach a system whose internals were opaque. Empowerment adds only the controllability facet, whether the recovered interface leaves the agent’s economy anything to decide, and on its own it is not a measure of agency (Sec. 9). Two caveats bound even this. Because the model is fit to uniform-action rollouts, what we recover is the world’s structure, the ideal belief an accurate agent would hold, not any particular agent’s posterior; and we exercise the controllability probe rather than the full three-criteria assessment, which the recovered model enables but which we leave to a study with agent-specific data.

Limitations. Results are on two small POMDPs with known ground truth, so validation is against an analytic model rather than blind; a state-count criterion stated before inspection, and a held-out protocol on data an exact model cannot memorise, would strengthen the claim. The model here is fit to exhaustive, uniform-action rollouts, so it recovers the environment’s generative structure; phenotyping a *particular* agent means fitting to that agent’s own rollouts, where the recovered model reflects its policy and where coverage of the state space is no longer guaranteed. Empowerment measures controllability, not agency: reading control regimes as agency profiles is licensed only because the agentic action chain is independently inspectable in the recovered model [1]. Finally, the description-length criterion for state count is conservative under our naive parameter count, which ignores gauge freedom; a gauge-aware count is future work.

Governance. Because the readout is architecture-agnostic and needs only a learnable model of the sensorimotor stream, it offers a concrete route to auditing opaque systems: fit an interpretable surrogate, recover its states and its controllability profile, and test whether governance handles track the recovered regimes [1]. Whether real interventions map onto these regimes is the empirical programme this pipeline makes testable.

11 Conclusion

Computational phenotyping of agency can be carried out end-to-end: from behaviour to a learned generative model, from that model to its recovered structure, and from the structure to a controllability phenotype that reproduces ground truth. The tensor network supplies exact conditionals and, once gauge-invariant states are fixed, a fully interpretable factored model—the inspectable surrogate that phenotyping of opaque systems requires.

Data and code. All code, models, and figures accompany the paper at <https://github.com/philipollaswilson/Tensor-Networks-for-the-Tmaze>.

Acknowledgments. AC was supported by ERC Synergy Grant (XScape) ERC-2020-SyG 951631.

Disclosure of Interests. The authors declare no competing interests.

References

1. Wilson, P., Albarracín, M., Hinrichs, N., Moore, J., Polani, D., Friston, K.J., Constant, A.: Characterizing Agency in AI Systems: Active Inference and an Empowerment-Based Probe (2025)
2. Parr, T., Pezzulo, G., Friston, K.J.: Active Inference. MIT Press (2022)
3. Klyubin, A.S., Polani, D., Nehaniv, C.L.: All else being equal be empowered. In: ECAL, pp. 744–753 (2005)
4. Salge, C., Glackin, C., Polani, D.: Empowerment—An Introduction. In: Guided Self-Organization, pp. 67–114. Springer (2014). https://doi.org/10.1007/978-3-642-53734-9_4
5. Montague, P.R., Dolan, R.J., Friston, K.J., Dayan, P.: Computational psychiatry. *Trends Cogn. Sci.* **16**(1), 72–80 (2012)
6. Schwartenbeck, P., Friston, K.: Computational phenotyping in psychiatry. *eNeuro* **3**(4) (2016). <https://doi.org/10.1523/ENEURO.0049-16.2016>
7. Wauthier, S.T., Vanhecke, B., Verbelen, T., Dhoedt, B.: Learning Generative Models for Active Inference Using Tensor Networks. In: IWAI. CCIS 1721, pp. 285–297. Springer (2023). https://doi.org/10.1007/978-3-031-28719-0_20
8. Heins, C., et al.: pymdp: A Python library for active inference in discrete state spaces. *J. Open Source Softw.* **7**(73), 4098 (2022)
9. Han, Z.-Y., Wang, J., Fan, H., Wang, L., Zhang, P.: Unsupervised Generative Modeling Using Matrix Product States. *Phys. Rev. X* **8**, 031012 (2018). [arXiv:1709.01662](https://arxiv.org/abs/1709.01662)
10. Glasser, I., Sweke, R., Pancotti, N., Eisert, J., Cirac, J.I.: Expressive power of tensor-network factorizations for probabilistic modeling. In: NeurIPS 32 (2019). [arXiv:1907.03741](https://arxiv.org/abs/1907.03741)
11. Hsu, D., Kakade, S.M., Zhang, T.: A spectral algorithm for learning hidden Markov models. *J. Comput. Syst. Sci.* **78**(5), 1460–1480 (2012). [arXiv:0811.4413](https://arxiv.org/abs/0811.4413)
12. Adhikary, S., Srinivasan, S., Miller, J., Rabusseau, G., Boots, B.: Quantum Tensor Networks, Stochastic Processes, and Weighted Automata. In: AISTATS. PMLR, vol. 130, pp. 2080–2088 (2021). [arXiv:2010.10653](https://arxiv.org/abs/2010.10653)

13. Chow, C.K., Liu, C.N.: Approximating discrete probability distributions with dependence trees. *IEEE Trans. Inf. Theory* **14**(3), 462–467 (1968)
14. Rissler, J., Noack, R.M., White, S.R.: Measuring orbital interaction using quantum information theory. *Chem. Phys.* **323**, 519–531 (2006). [arXiv:cond-mat/0508524](#)
15. Convy, I., Huggins, W., Liao, H., Whaley, K.B.: Mutual information scaling for tensor network machine learning. *Mach. Learn.: Sci. Technol.* **3**, 015017 (2022). [arXiv:2103.00105](#)
16. Friston, K.J., Parr, T., Zeidman, P.: Bayesian model reduction. [arXiv:1805.07092](#) (2019)
17. Smith, R., Schwartenbeck, P., Parr, T., Friston, K.J.: An active inference approach to modeling structure learning. *Front. Comput. Neurosci.* **14**, 41 (2020)
18. Orús, R.: A practical introduction to tensor networks: MPS and PEPS. *Ann. Phys.* **349**, 117–158 (2014). [arXiv:1306.2164](#)
19. Aizpurua, B., Palmer, S., Orús, R.: Tensor Networks for Explainable Machine Learning in Cybersecurity. [arXiv:2401.00867](#) (2024)
20. Li, C., Zeng, J., Tao, Z., Zhao, Q.: Permutation Search of Tensor Network Structures via Local Sampling. In: *ICML (2022)*. [arXiv:2206.06597](#)
21. Li, C., Zeng, J., Li, C., et al.: Alternating Local Enumeration (TnALE). In: *ICML (2023)*. [arXiv:2304.12875](#)